

# Solver in VeGen

The role of the solver in the VeGen framework was initially unclear to me. I discovered that the `optimizeBottomUp` function, which is invoked in the GSLP driver, is the core of the solver. The solver aims to create a good `VectorPackSet` for later code generation, using planning and heuristics. The construction of a `VectorPackSet` is simple, but the codegen method is complex. The only interaction between the solver and the `VectorPackSet` is the `tryAdd` method, which checks the compatibility of the `VectorPack` to be added and updates the `PackedValues`, `AllPacks`, and `ValueToPackMap`.

The main logic of the solver is implemented in another `optimizeBottomUp` function, which is overloaded with `Packer` and `SeedOperands` as parameters and returns a vector of `VectorPacks` to form a `VectorPackSet`. The solver is not a class, but a function named `optimizeBottomUp` in the solver file.

## 1 `optimizeBottomUp`

The function `optimizeBottomUp` is a solver that tries to find the best vectorization plan for a given set of seed operands. It takes four parameters:

- `Packs`: a reference to a vector of `VectorPack` pointers, which are data structures that represent a group of scalar instructions that can be packed into a single vector instruction. The function will append the optimal vector packs to this vector.
- `Pkr`: a pointer to a `Packer` object, which is a class that handles the packing of vector instructions. It also provides access to the LLVM context and the target machine information.
- `SeedOperands`: an array of `OperandPack` pointers, which are data structures that represent a group of scalar operands that can be packed into a single vector operand. These are the starting points for the vectorization process.
- `BlocksToIgnore`: blocks that should be skipped.

The function performs the following steps:

- It creates a `CandidatePackSet` object, which is a data structure that stores a set of candidate vector packs and a mapping from scalar instructions to vector packs that contain them.
- It resizes the `Inst2Packs` vector, which is a member of the `CandidatePackSet` object, to match the number of values in the `VectorPackContext`. This vector stores a list of vector packs for each scalar instruction.
- It iterates over the candidate vector packs and updates the `Inst2Packs` vector accordingly. For each vector pack, it gets the bitset of the elements that are packed, and for each set bit, it adds the vector pack to the corresponding list in the `Inst2Packs` vector.
- It creates a `Plan` object, which represents a vectorization plan. It takes the packer as a parameter and initializes the plan with the scalar cost, which is the total cost of executing the scalar instructions without vectorization.
- It calls the `improvePlan` function, which is defined in the same file as `optimizeBottomUp`, to find the optimal vectorization plan for the given seed operands and candidate vector packs. This function uses a greedy algorithm that tries to reduce the cost of the plan by adding or removing vector packs, while respecting the dependencies and constraints of the vectorization process.
- It inserts the vector packs from the optimal plan to the `Packs` vector, which is the output parameter of the function.

- It calls the `findDepCycle` function, which is defined in the same file as `optimizeBottomUp`, to check if the vectorization plan introduces any dependence cycles that would prevent the correct execution of the program. If so, it prints an error message, clears the `Packs` vector, and returns the scalar cost as the result of the function.
- Otherwise, it returns the cost of the optimal vectorization plan as the result of the function.

## 2 Plan

A `VPlan` is a data structure used by `VeGen`, a vectorization framework for LLVM. A `VPlan` contains a reference to a `Packer` object, a floating-point number `Cost`, and a set of `VectorPack` objects. As the name suggests, a `VectorPackSet` is also a collection of `VectorPack` objects, but a `VPlan` is more concerned with the cost and the search algorithm.

## 3 improvePlan

The function name is somehow misleading, since originally there is no plan at all. It is `improvePlan` that plans from the beginning.

### 3.1 Seeds

The `improvePlan` function takes a vector of `VectorPack` objects as `Seeds`.

#### 3.1.1 Store Packs

- It iterates over the instructions in the function obtained from the input `Packer`, and ignores those that are not scalar store instructions.
- Then, it calls `getSeedMemPacks`, which uses `AccessDAG` to pack consecutive scalar stores into a `VectorPack`. When generating `VectorPacks`, the function sets the vector length (`VL`) to 2, 4, or 8, and adds the generated packs to the `Seeds`.
- Finally, it sorts the `VectorPacks` in `Seeds` according to the `id` in `LayoutInfo` and the descending order of their element counts, which prioritizes smaller `ids` (i.e., the leaders in access) and larger element counts.

#### 3.1.2 Loop-Reduction Packs

A loop reduction variable is updated in every loop iteration, so it must be compiled into a  $\phi$  instruction in the SSA IR. Starting from the  $\phi$  instructions, the `improvePlan` function recursively builds a reduction chain, producing a `ReductionInfo` object. The reduction variable can be either integer or floating-point, and it cannot have more than one use. Moreover, the reduced values that form the reduction variable cannot have more than one in-loop use (otherwise, after vectorizing, we have to extract the value from the vector for that use). The function also takes special care of the operands of the call instructions, trying to find the reduction opportunities of these operands.

#### 3.1.3 Heuristic

The `Heuristic` module in `VeGen` is responsible for finding a suitable solution for a given operand pack. The interface of `Heuristic` has only one important function, namely `solve`, which takes a constant reference to an `OperandPack` and returns a `Solution` object. A `Solution` object contains an array of `VectorPacks` that represent the vectorized code (similar to a `VPlan`).

The Heuristic module is based on three simple rules, which correspond to three heuristic ways of building vector packs. They are: (1) insert all operands into a vector pack, (2) broadcast an operand if all elements in a vector pack are equal, and (3) extract an operand from an existing vector pack.

An OperandPack must be a vector, not a set, because one Value could be used multiple times. To enable more flexibility in using the extracted values, the Heuristic module also allows two shuffle strategies, transpose and de-interleave, to rearrange the elements in a vector pack.

In the Solver module, the Heuristic module has two uses. One is inside the improvePlan function, which is passed to the runBottomUpFromOperand function. The other is inside the tryPackBackEdgeConditions function, which is used to handle loop-carried dependencies.

### 3.1.4 runBottomUpFromOperand

runBottomUpFromOperand takes five parameters: a pointer to an OperandPack object, a reference to a Plan object, a reference to a Heuristic object, a boolean value, and a function object. The purpose of this function is to optimize a plan for vectorizing a set of operands using a bottom-up heuristic.

An OperandPack is a data structure that represents a group of scalar operands that can be vectorized together. A VectorPack is a data structure that represents a group of vector operands that are produced by vectorizing an OperandPack. A Plan is a data structure that stores a set of VectorPacks and their associated costs. A Heuristic is a data structure that implements a strategy for finding the best VectorPacks for a given OperandPack.

The function body consists of the following steps:

- I. Initialize a Worklist of OperandPack pointers, and push the input OperandPack pointer to it. This Worklist will store the OperandPacks that need to be processed by the heuristic.
- II. Initialize a Visited set of OperandPack pointers. This Visited will store the OperandPacks that have been already processed by the heuristic.
- III. Loop until the Worklist is empty:
  - i. Verify the cost of the current Plan using the verifyCost method.
  - ii. Pop the last OperandPack pointer from the WorkList and assign it to a local variable named OP.
  - iii. If OP is already in the Visited set, skip the rest of the loop iteration. This avoids processing the same OperandPack twice.
  - iv. Call the solve method of the Heuristic object with OP as the argument, and assign the result to a local variable named NewPacks. This is a SmallVector of VectorPack pointers that represents the best solution for vectorizing OP according to the heuristic.
  - v. Initialize a local variable named OldPacks as a SmallPtrSet of VectorPack Pointers. This will store the VectorPacks that need to be replaced by the NewPacks.
  - vi. Loop over the NewPacks vector:
    - a) Loop over the element values of each VectorPack using the elementValues method. There are the scalar values that are combined into the vector value.
    - b) If the element value is an Instruction object, call the getProducer method of the Plan object with it as the argument, and assign the result to a local variable named VP2. This is a VectorPack pointer that represents the current producer of the element value in the Plan.

- c) If VP2 is not null, insert it into the OldPacks set. This means that the element value is already vectorized by another VectorPack that needs to be replaced by the new one.
- vii. If the OverrideExisting parameter is false and the OldPacks set is not empty, skip the rest of the loop iteration. This means that the function does not allow replacing existing VectorPacks in the Plan.
- viii. Loop Over the OldPacks set:
  - a) Call the remove method of the Plan object with each VectorPack pointer as the argument. This removes the VectorPack from the Plan and updates its cost.
- ix. Loop over the NewPacks vector:
  - a) Call the add method of the Plan object with each VectorPack pointer as the argument. This adds the VectorPack to the Plan and updates its cost.
  - b) Call the getOperandPacks method of each VectorPack and assign the result to a local variable named Operands.
  - c) Append the Operands array to the end of the WorklistVector using the append method. This adds the operands of the VectorPack to the Worklist for further processing by the heuristic.
  - d) If the GetExtraOperands parameter is not null, call it with the VectorPack pointer and the Worklist vector as the arguments. This is a function object that can add extra operands to the Worklist based on custom logic.
- IV. The function does not return anything, but it modifies the Plan object passed by reference. The final Plan object represents the optimized plan for vectorizing the input OperandPack and its descendants using the bottom-up heuristic.

### 3.1.5 Decomposed Stores

decomposeStorePacks takes two parameters: a pointer to a Packer object and a pointer to a VectorPack object. The purpose of this function is to decompose a store pack into smaller packs that fit in the native bitwidth of the target architecture.

A Packer is a data structure that implements the vectorization algorithm for a given function. A VectorPack is a data structure that represents a group of vector operands that are produced by vectorizing an OperandPack. An OperandPack represents a group of scalar operands that can be vectorized together. A StoreInst is a subclass of Instruction that represents a store operation.

- Assert that the input VectorPack is a store pack using the isStore method. This means that it contains store instructions as its element values.
- Call the getOrderedValues method of the input VectorPack and assign the result to a local variable named Values. This is an ArrayRef of Value pointers that represents the ordered element values of the store pack.
- Cast the first element of Values to a StoreInst, SI.
- Call the getLoadStoreVecRegBitWidth method of the TTI object, and divide the result by the bit width of the value operand of the SI using the getBitWidth function and the getDataLayout method of the Packer object. Assign the result to a local variable named VL. This is an unsigned integer that represents the maximum vector length in bits that can be used for load and store operations.

- Call the `getDA` method of the `Packer` to get a `DependenceAnalysis` result, `DA`.
- if the size of the `Values` array is less than or equal to the `VL` variable, return a `SmallVector` containing the input `VectorPack` as the only element. This means that the store pack does not need to be decomposed further.
- Otherwise, initialize a local variable named `Decomposed` as a `SmallVector` of `VectorPack` pointers. This will store the decomposed store packs.
- Loop over the `Values` array with an index variable named `i` that starts from zero and increments by `VL` until it reaches the size of the `Values` array, which is assigned to a local variable named `N`. This loop will split the store pack into smaller chunks of size `VL`.

### 3.1.6 `tryPackBackEdgeConds`

A `conditionPack` represents a group of control conditions that affect the vectorization of an `OperandPack`. A `Heuristic` implements a strategy for finding the best `VectorPacks` for a given `OperandPack`.

The `getBackEdgePacks` function takes three parameters: a pointer to a `Packer`, a reference to a `Plan`, and a reference to a `SmallPtrSet` of `ConditionPack` pointers. The purpose of this function is to collect all the back-edge condition packs for packs of values from divergent loops in the plan.

A back-edge condition is a control condition that determines whether a loop will iterate again or exit. A divergent loop is a loop that has different trip counts for different vector lanes. A trip count is the number of times a loop iterates. A vector lane is a scalar element of a vector value.